

CSC-496/696: Natural Language Processing and Text as Data

Lecture 15: Transformers

Patrick Wu

Friday, October 25, 2024

Lecture Contents

1. Announcements
2. Review of the Attention Mechanism
3. Attention is All You Need
4. Transformers

Announcements

Midterm

- Midterm solutions are now uploaded
- I forgot to upload it last week...sorry.

Final Project

- A more detailed write-up of the final project is posted on Canvas
- Also includes information about the dataset for the task-driven version of the final project, if that is what you're interested in
- **I strongly suggest coming to office hours to discuss project ideas or approaches to the dataset that has been posted**

Final Project Data Posted

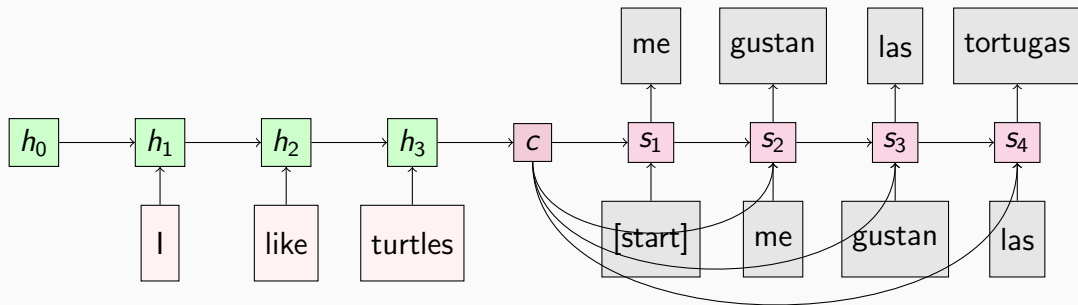
- The data for the task-driven version of the project is also posted
- You are also free to use this data for whatever research question you want
- Remember, the research question does not have to be groundbreaking. It doesn't even have to be novel—there is utility in replicating results using different approaches.

Assignment 3

- Due on Tuesday, October 29
- Future assignments will be less work to make time to complete the final project

Review of the Attention Mechanism

Encoder-Decoder with Context Sharing



$$h_t = f_W(x_t, h_{t-1})$$
$$s_t = g_U(y_{t-1}, s_{t-1}, c)$$

- Why not just average h_1 , h_2 , and h_3 ?

- Why not just average h_1 , h_2 , and h_3 ?
 - Effectively the same as c
 - We want to emphasize certain encoder hidden states over others at each decoding step
 - In other words, we want a dynamic context vector c_t for each decoder step

- Why not just average h_1 , h_2 , and h_3 ?
 - Effectively the same as c
 - We want to emphasize certain encoder hidden states over others at each decoding step
 - In other words, we want a dynamic context vector c_t for each decoder step
- We use the *attention mechanism* to operationalize this idea

Intuition of Attention

- When you're reading, you connect previous information to what you're currently reading, but you don't weigh all previous information in the same way
- You are paying different amounts of *attention* to the previous information
- Your attention to previous information can also shift as you read further along

Intuition of Attention

- Similarly, the attention mechanism learns which encoder hidden states need more “attention” (weighed more) and which encoder hidden states need less “attention” (weighed less)
- Dot-product attention uses the dot product to score how closely related a decoder hidden state is to every encoder hidden state
- We'll use the dot product to calculate a score between s_{t-1} and all encoder hidden states, h_j

$$\text{score}(s_{t-1}, h_j) = s_{t-1} \cdot h_j$$

Dot-Product Attention

After calculating this score for all encoder hidden states, we can then calculate the weights using the softmax

$$\alpha_{tj} = \text{softmax}(\text{score}(s_{t-1}, h_j)) = \frac{\exp(\text{score}(s_{t-1}, h_j))}{\sum_k \exp(\text{score}(s_{t-1}, h_k))}$$

α_{tj} is the weight between the decoder hidden state s_{t-1} and encoder hidden state h_j

Dot-Product Attention

Finally,

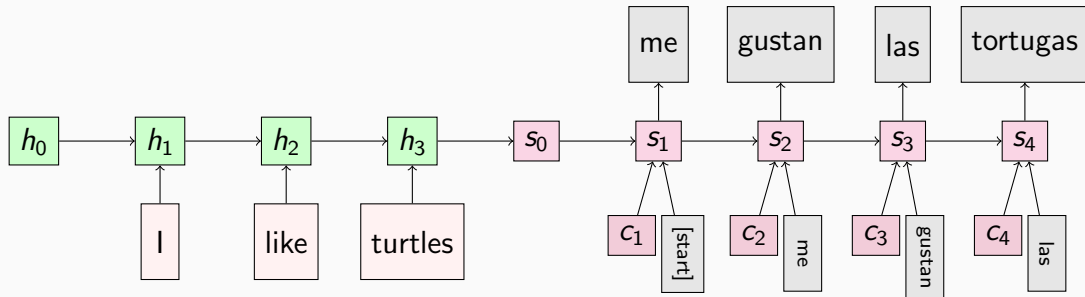
$$c_t = \sum_k \alpha_{tk} h_k$$

We can then calculate s_t as follows

$$s_t = g_U(y_{t-1}, s_{t-1}, c_t)$$

Notice that c_t is a weighted sum of *all* of the hidden state of the encoder. By backpropping through everything, we can learn which encoder states to pay more “attention” to at each decoding step.

Encoder-Decoder with Dot-Product Attention



Example: Calculating c_2

c_2 is the weighted sum of the encoder hidden states, h_1 , h_2 , and h_3 . Note that the initial hidden state h_0 is **not** included in the calculation.

$$c_2 = \sum_{k=1}^3 \alpha_{2k} h_k$$

We need to calculate α_{21} , α_{22} , and α_{23} .

Example: Calculating c_2

We calculate the dot products between s_1 (the previous hidden state of the decoder) and all the hidden states of the encoders.

$$\text{score}(s_1, h_1) = s_1 \cdot h_1 = \gamma_1$$

$$\text{score}(s_1, h_2) = s_1 \cdot h_2 = \gamma_2$$

$$\text{score}(s_1, h_3) = s_1 \cdot h_3 = \gamma_3$$

We can put these dot products into a vector:

$$\begin{bmatrix} \gamma_1 & \gamma_2 & \gamma_3 \end{bmatrix}$$

Example: Calculating c_2

We α values are calculated using the softmax over this vector

$$\alpha = \left[\frac{\exp(\gamma_1)}{\exp(\gamma_1) + \exp(\gamma_2) + \exp(\gamma_3)} \quad \frac{\exp(\gamma_2)}{\exp(\gamma_1) + \exp(\gamma_2) + \exp(\gamma_3)} \quad \frac{\exp(\gamma_3)}{\exp(\gamma_1) + \exp(\gamma_2) + \exp(\gamma_3)} \right]$$

Notice that this respects the rules of the softmax we had previously discussed: all numbers are positive, between 0 and 1, and all three numbers add up to 1.

The first element is α_{21} , the second element is α_{22} , and the third element is α_{23} .

Example: Calculating c_2

Then,

$$c_2 = \alpha_{21}h_1 + \alpha_{22}h_2 + \alpha_{23}h_3$$

Notice that h_1 , h_2 , and h_3 are vectors, so c_2 is a vector with the dimension of the encoder hidden states. In other words, c_2 **is a weighted sum of the encoder hidden states**.

c_2 is then used with the input “me” and s_1 in order to calculate s_2 .

- When you read papers, you might run across the term “attending”

Attending

- When you read papers, you might run across the term “attending”
- This is just the mechanism by which each decoder state pays a different amount of “attention” (α weights) to each encoder hidden state

Attending

- When you read papers, you might run across the term “attending”
- This is just the mechanism by which each decoder state pays a different amount of “attention” (α weights) to each encoder hidden state
- Notice also that we attend to *every* encoder hidden state

Attending

- When you read papers, you might run across the term “attending”
- This is just the mechanism by which each decoder state pays a different amount of “attention” (α weights) to each encoder hidden state
- Notice also that we attend to *every* encoder hidden state
- No matter how inconsequential an encoder hidden state is, it will have *some* weight in the weighted average

Attending

- When you read papers, you might run across the term “attending”
- This is just the mechanism by which each decoder state pays a different amount of “attention” (α weights) to each encoder hidden state
- Notice also that we attend to *every* encoder hidden state
- No matter how inconsequential an encoder hidden state is, it will have *some* weight in the weighted average
- Called soft attention

Soft vs. Hard Attention

- Soft attention: attends to every encoder hidden state, even if the weight is very low
 - Advantage: differentiable, easy to train using backprop
 - Disadvantage: computationally more expensive
- Hard attention: focuses on specific parts, attending only to a specific subset of encoder hidden states
 - Advantage: more computationally efficient
 - Disadvantage: non-differentiable, which means we require reinforcement learning methods

Review Question 1

The dynamic context vector for each step of the decoder is a scalar value (a single number)

- (A) True
- (B) False

Review Question 2

An attention weight is a scalar value (a single number)

(A) True

(B) False

Review Question 3

What is the dimension of a context vector?


- (A) The dimension of the hidden states of the decoder
- (B) The dimension of the input embedding of the encoder
- (C) The dimension of the output embedding of the encoder
- (D) The dimension of the hidden states of the encoder

Disadvantage of RNNs/LSTMs

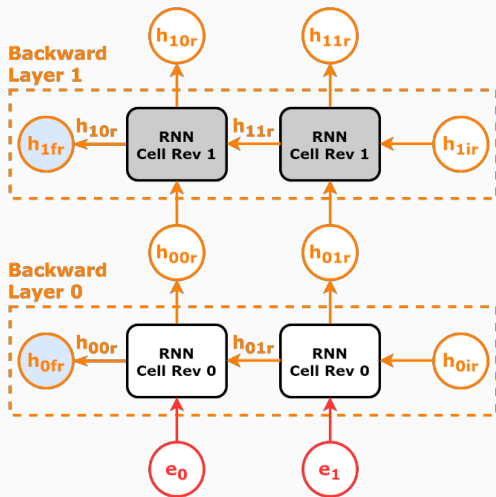
- Researchers found *significant* improvements when using attention
- Widely explored in contexts such as text summarization, image captioning, and machine translation
- However, computations are still inherently sequential
 - Still must feed in one input at a time
 - We can parallelize using multiple sequences (batching), but we can't parallelize within the sequence
 - For example, we cannot calculate h_3 without calculating h_2

- By no means are RNNs *bad*—still used for tasks like time-series predictions!
- In 2018, AI2 and the University of Washington released one of the first “large” language models called **E**MBEDDINGS FROM **L**ANGUAGE **M**ODEL, or ELMo
- ELMo is a model based on bidirectional LSTMs
- Because it was not an encoder-decoder framework, it **did not use attention**

Forward Language Model



Backward Language Model



Pretrained on Two Tasks

ELMo was trained on two tasks:

- Forward language modeling: ELMo predicts the next word in a sequence given the previous words

$$P(w_t | w_1, w_2, \dots, w_{t-1})$$

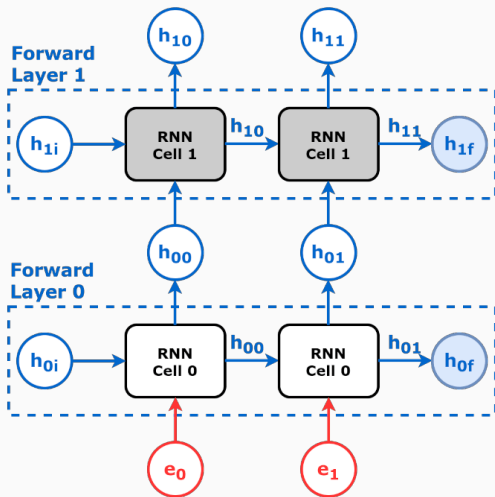
- Backward language modeling: ELMo processes the text in reverse, from right to left, predicting the previous word based on future words in the sequence

$$P(w_t | w_{t+1}, w_{t+2}, \dots, w_T)$$

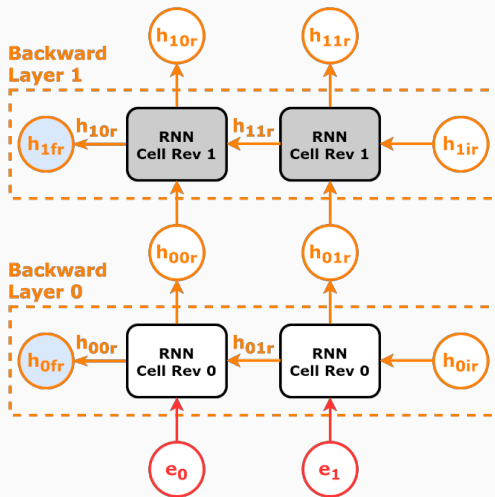
Contextual Word Embeddings

- ELMo produces contextual word embeddings
- That is, the word embedding for the word “bank” changes based on its context
- In stark contrast to the previous word embedding methods, which produced a single word embedding for each unique word in the corpus

Forward Language Model



Backward Language Model



Attention is All You Need

Background

- Proposed in paper, “Attention is All You Need” by Vaswani et al. (2017)
 - Paper can be found [here](#)
- The original transformer, as described in the paper, is an encoder-decoder model
- Note: depending on the book or paper, people might capitalize the T in “transformer.” Because transformers are so ubiquitous now in deep learning, most new papers do not capitalize the t.

Intuition of Transformers

- The original paper focused on machine translation
- Encoder module: processes the input text and encodes it into a series of vectors that capture the contextual information of the input
- Decoder module: takes these encoded vectors and generates the output text

Intuition of Transformers

- Much like an RNN, transformers take as input word embeddings
- A transformer **re-expresses** (“transforms”) each word embedding as a weighted sum of all word embeddings in its context
- Recall: when using dot-product attention with RNNs, we calculated a context vector that would be unique for each decoder step by using a weighted sum of the hidden states of the encoder
- We'll use another form of attention here: **self-attention**

Intuition of Transformers

- The self-attention mechanism allows the model to weigh the importance of different words or tokens in a sequence relative to each other
- Allows the model to capture long-range dependencies and contextual relationships within the input data
 - Addresses a key weakness of RNNs that was partially addressed by attention
- In other words, self-attention allows each position in the input sequence to “attend to” all other positions in the same sequence when calculating the representation of a sequence

Why Self-Attention?

- Consider the sentence: “The chicken didn’t cross the road because it was too tired.”

Why Self-Attention?

- Consider the sentence: “The chicken didn’t cross the road because it was too tired.”
- Consider another sentence: “The chicken didn’t cross the road because it was too wide.”

Why Self-Attention?

- Consider the sentence: “The chicken didn’t cross the road because it was too tired.”
- Consider another sentence: “The chicken didn’t cross the road because it was too wide.”
- “It” refers to different nouns in each sentence.
 - In the first sentence, “it” refers to the chicken.
 - In the second sentence, “it” refers to the road.

Why Self-Attention?

- Consider the sentence: “The chicken didn’t cross the road because it was too tired.”
- Consider another sentence: “The chicken didn’t cross the road because it was too wide.”
- “It” refers to different nouns in each sentence.
 - In the first sentence, “it” refers to the chicken.
 - In the second sentence, “it” refers to the road.
- If we read the sentences from left to right, we get: The chicken didn’t cross the road because it

Why Self-Attention?

- Consider the sentence: “The chicken didn’t cross the road because it was too tired.”
- Consider another sentence: “The chicken didn’t cross the road because it was too wide.”
- “It” refers to different nouns in each sentence.
 - In the first sentence, “it” refers to the chicken.
 - In the second sentence, “it” refers to the road.
- If we read the sentences from left to right, we get: The chicken didn’t cross the road because it
- At this point, we don’t know what “it” is referring to.

Why Self-Attention?

- Takeaway: language is not something that is *relationally* left to right
- The English language is read left to right, but words relate to each other in both directions
- One of the fundamental limitations of RNNs was that you must walk through the sequence one word at a time
- Attention had solved many of these issues, but again, we cannot calculate the context vectors c_t in parallel

Self-Attention (Simplified)

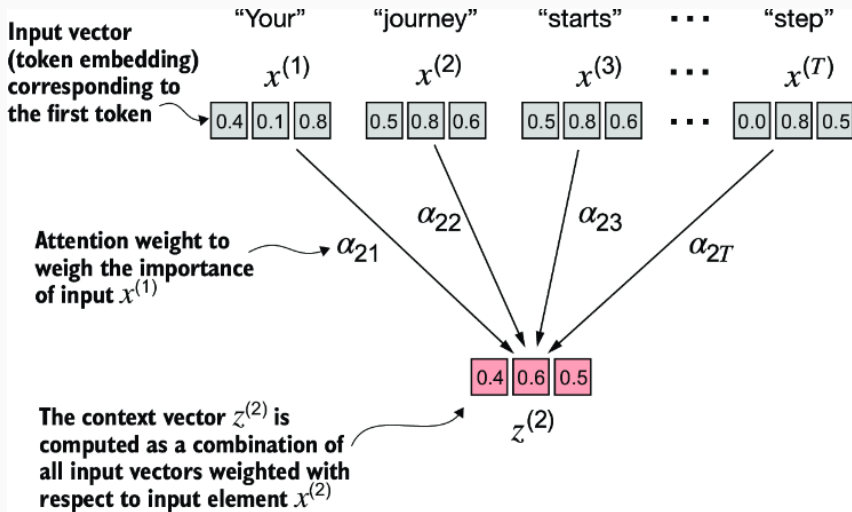


Figure from Raschka (2024). $z^{(2)}$ is calculated using attention weights with respect to

Self-Attention (Simplified)

Specifically (in this simplified version of self-attention),

$$z^{(2)} = \sum_{t=1}^T \alpha_{2t} x^{(t)}$$

where

$$\alpha_{2t} = \text{softmax}(\text{score}(x^{(2)}, x^{(t)}))$$

for all $t \in \{1, \dots, T\}$, and

$$\text{score}(x^{(2)}, x^{(t)}) = x^{(2)} \cdot x^{(t)}$$

Self-Attention (Simplified)

- This is why it's called *self*-attention—it is the attention the sentence “pays to itself”

Self-Attention (Simplified)

- This is why it's called *self*-attention—it is the attention the sentence “pays to itself”
- We end up with features $z^{(1)}, z^{(2)}, \dots, z^{(T)}$ that are transformed features of the input features

Self-Attention (Simplified)

- This is why it's called *self*-attention—it is the attention the sentence “pays to itself”
- We end up with features $z^{(1)}, z^{(2)}, \dots, z^{(T)}$ that are transformed features of the input features
- Each transformed feature is a weighted sum of all the other features

Self-Attention (Simplified)

- This is why it's called *self*-attention—it is the attention the sentence “pays to itself”
- We end up with features $z^{(1)}, z^{(2)}, \dots, z^{(T)}$ that are transformed features of the input features
- Each transformed feature is a weighted sum of all the other features
- Model can learn *very* complex relationships between the input features!

Self-Attention: More Details

In self-attention, each input vector plays three roles:

- Query: the current element being compared to all other inputs
 - In the previous example, $x^{(2)}$ is the query
- Key: the input being compared to the query to determine the similarity weight
 - In the previous example, this is the x_t in $\text{score}(x^{(2)}, x^{(t)}) = x^{(2)} \cdot x^{(t)}$
- Value: the input that gets weighted and summed up to compute the output for the current element
 - In the previous example, recall that $z^{(2)} = \sum_{t=1}^T \alpha_{2t} x^{(t)}$. $x^{(t)}$ is called the value.

Self-Attention (Full Version)

Because each vector plays three roles, we can add three weight matrices that will dramatically increase the ability to capture even more nuanced patterns

$$q^{(i)} = x^{(i)} W^q$$

$$k^{(i)} = x^{(i)} W^k$$

$$v^{(i)} = x^{(i)} W^v$$

Eventually, we will learn each weight matrix through backpropagation.

Self-Attention (Full Version)

$$q^{(i)} = x^{(i)} W^q; k^{(i)} = x^{(i)} W^k; v^{(i)} = x^{(i)} W^v$$

$$\text{score}(x^{(i)}, x^{(t)}) = \frac{q_i \cdot k_t}{\sqrt{d_k}}$$

$$\alpha_{it} = \text{softmax}(\text{score}(x^{(i)}, x^{(t)})) \quad \forall t \in \{1, \dots, T\}$$

$$z^{(i)} = \sum_{t=1}^T \alpha_{it} x^{(t)}$$

Self-Attention (Full Version)

- Notice that we calculate the score by using the dot product, but also dividing by the square root of the key
- This is called scaled dot-product attention
- Prevents the dot product from blowing up
- But unlike using cosine similarity, it doesn't limit the dot product to between -1 and 1


Review Question 4

What is the primary function of self-attention?

- (A) To attend to different parts of the input sequence by focusing on its own elements
- (B) To combine input sequences with random weights
- (C) To ignore irrelevant information from the input and focus on a fixed part
- (D) To compute the overall sentiment of the sentence

Multi-head Self-Attention

- What we just described is a **single attention head**
- In transformers, we can use multiple heads; we can even greater diversity in representing the inputs because the weight matrices associated with the query, key, and value can all be different
- Intuition: multi-head attention allows each self-attention head to learn distinct patterns within the sentence
- We can concatenate the output of each attention head, and then project it down to our original input dimensionality using a weight matrix W^O
- Equations (9.14) to (9.19) in SLP Ch. 9 shows all the mathematical details of multi-headed attention



Unmasked vs. Masked Attention

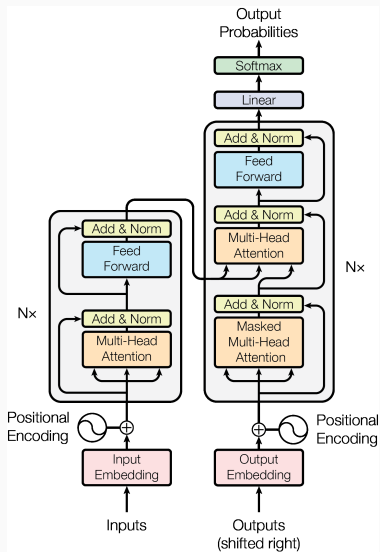
- In the above, we allowed each query to attend to all preceding *and* subsequent outputs
- But sometimes we don't want to do that—we want the model to only learn from the preceding input
 - Especially important if we are training generative models!!
- We can use masked attention, where the model is only allowed to attend to previous inputs
- Using our running example, $z^{(2)}$ would only be calculated using $x^{(2)}$ and $x^{(1)}$

Transformers


Self-Attention and Transformers

- The paper is called “Attention is All You Need” because the authors show that you only need the self-attention mechanism to process sequential data
- By dispensing with recurrence, we can parallelize calculations, yielding much faster training time
- Self-attention is the heart of transformers

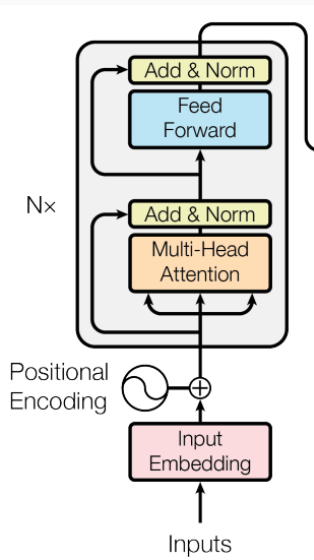
Transformers (Full Diagram)



Encoder-Decoder Architecture of the Transformer



Transformer Encoder



Encoder: Multi-Head Self-Attention

This is just what we talked about in the first half of this lecture

Encoder: Feed Forward

- The feed-forward sublayer in the encoder (and the decoder) is a simple two-layer fully connected neural network
- But instead of processing the whole sequence of embeddings as a single vector, it processes each embedding independently
- The hypothesis is that this is where most of the capacity and memorization happens
- Typically, transforms the input features from k elements to $4k$ elements, and then back down to k elements

Encoder: Add & Norm

- A transformer uses skip (or residual) connections
 - Won't talk about the details in this class, but it is a concept borrowed from computer vision, which allowed convolutional neural networks to be rapidly scaled up
- Layer normalization normalizes each input in the batch to have zero mean and a variance of 1
- Different models move around the layer normalization for training efficiency purposes

Encoder: Positional Encoding

- Recall that one issue is that unmasked self-attention does not recognize word order—it just calculates the relevancy between a given query and the keys

Encoder: Positional Encoding

- Recall that one issue is that unmasked self-attention does not recognize word order—it just calculates the relevancy between a given query and the keys
- We're going to add more embeddings: augment the token embeddings with a position-dependent vector

Encoder: Positional Encoding

- Recall that one issue is that unmasked self-attention does not recognize word order—it just calculates the relevancy between a given query and the keys
- We're going to add more embeddings: augment the token embeddings with a position-dependent vector
- These embeddings can learn information about the positions of words

Encoder: Positional Encoding

- Recall that one issue is that unmasked self-attention does not recognize word order—it just calculates the relevancy between a given query and the keys
- We're going to add more embeddings: augment the token embeddings with a position-dependent vector
- These embeddings can learn information about the positions of words
- In fancier terms, self-attention and the feed-forward layers are said to be permutation equivariant

Encoder: Positional Encoding


- Recall that one issue is that unmasked self-attention does not recognize word order—it just calculates the relevancy between a given query and the keys
- We're going to add more embeddings: augment the token embeddings with a position-dependent vector
- These embeddings can learn information about the positions of words
- In fancier terms, self-attention and the feed-forward layers are said to be permutation equivariant
- Positional embeddings break this permutation equivariance

Encoder: Positional Encoding


- Recall that one issue is that unmasked self-attention does not recognize word order—it just calculates the relevancy between a given query and the keys
- We're going to add more embeddings: augment the token embeddings with a position-dependent vector
- These embeddings can learn information about the positions of words
- In fancier terms, self-attention and the feed-forward layers are said to be permutation equivariant
- Positional embeddings break this permutation equivariance
- There are a few ways to calculate or learn these position embeddings, but we won't talk about these in this class

Stacking Encoders

- You can also *stack* encoders
- That is, you input your token embeddings into an encoder, go through the entire transformer encoder *block*, and then feed the output of that transformer encoder block into another transformer encoder block



Transformer Decoder



Decoder: Masked Multi-Head Self-Attention

- Ensures that the tokens we generate at each time step are only based on the past outputs and the current token being predicted
- We don't want the decoder to cheat by “peeking” ahead
- In other words, the current input can only attend to the previous inputs and itself, but not the subsequent inputs

Decoder: Encoder-Decoder Attention Layer

- Performs multi-head attention over the output values of the encoder stack, which act as the keys and values
- The intermediate representations of the decoder act as the queries
- This is the part of the decoder that relates the encoder output with the decoder input
 - Where language translation occurs!

Stacking Decoders

- You can also stack decoders as well
- Operates in a similar logic as stacking encoders

- And that's it! We've covered transformers!
- Even though it looks complicated, we've now broken down (conceptually, at least) all the details of both the encoder and decoder
- Original paper focused on machine translation
- But as the paper captured the imaginations of many researchers, people realized that you could model language using only encoders or decoders

Large Language Models

- Modern large language models (LLMs) are all based on transformers
- There are encoder-only models (e.g., BERT), decoder-only models (e.g., GPT), and encoder-decoder models (e.g., T5)
- Decoder-only models are currently the center of attention at this moment, but all approaches dramatically improved benchmark results
- Large language models also allow us to *pretrain* a model on huge amounts of data, which then allows us to either use them directly for tasks of interest or we can *finetune* a model with a small amount of labeled data